

RAMANJANEYULU SIMHADRI

Cincinnati, OH | +1 (217) 691-3997 | rsimha084@gmail.com | [LinkedIn](#) | [Portfolio](#)

PROFESSIONAL SUMMARY

- Data Engineer with 4+ years of experience building ETL pipelines and analytics datasets using SQL, Snowflake, dbt, Python, and Spark for financial reporting and customer analytics.
- Developed batch and streaming data pipelines using Azure Data Factory, Synapse Analytics, ADLS, and Kafka, improving ETL processing efficiency and reporting availability by up to 30%.
- Optimized Snowflake and Redshift warehouse transformations using SQL tuning and dimensional modeling, improving dataset reliability for Business Intelligence and segmentation reporting.
- Collaborated with analytics and ingestion teams to deliver governed transformation-layer datasets aligned with DataOps practices supporting enterprise analytics and Data Platform Engineer workflows.

TECHNICAL SKILLS

Data Engineering & ETL: Data Engineering, Data Pipelines, ETL Development, Data Transformation, Data Integration, Data Modeling (Dimensional & Relational), Analytics Engineering.

Programming & Query Languages: SQL, Python, PySpark.

Data Warehousing & Big Data: Snowflake, Amazon Redshift, Apache Spark, Kafka.

Cloud Data Platforms: AWS (S3, Glue, Redshift), Azure (Data Factory, Databricks, Synapse Analytics, ADLS).

Pipeline Orchestration: dbt, Apache Airflow, Azure Data Factory.

Database Engineering: Relational Databases, SQL Optimization, Query Performance Tuning.

Analytics & Business Intelligence: Business Intelligence Data Modeling, Reporting Dataset Engineering.

Data Quality & Governance: Data Validation Frameworks, Data Lineage, Metadata Management.

Monitoring & Observability: CloudWatch, Azure Monitor, Splunk.

Engineering Practices: Git, Agile, DataOps Collaboration.

PROFESSIONAL EXPERIENCE

Data Engineer

Aug 2023 - Present

Capital One | McLean, VA

- Designed Snowflake SQL transformation pipelines implementing customer householding relationship logic across multi-million record datasets, improving segmentation dataset readiness for analytics teams.
- Built modular dbt transformation models with automated testing and documentation layers, reducing manual validation effort during production refresh cycles by 30%.
- Optimized complex join strategies and aggregation logic in Snowflake using clustering-aware query tuning, decreasing transformation runtime by 28% across monthly refresh workloads.
- Collaborated with Business Intelligence analysts to structure dimensional reporting datasets supporting customer growth measurement and campaign-performance tracking dashboards.
- Implemented relational transformation-layer validation checks using SQL data-quality controls that reduced downstream reconciliation issues during scheduled dataset releases.
- Resolved upstream schema inconsistencies by coordinating with ingestion engineering teams, stabilizing transformation dependencies across enterprise refresh pipelines.
- Improved transformation lineage transparency by documenting dbt model dependencies and model relationship mappings, reducing troubleshooting turnaround time by approximately 25%.
- Strengthened reporting dataset usability by refining dimensional modeling structures supporting segmentation workflows used across multiple analytics stakeholder groups.

Data Pipeline Engineer

May 2021 - Dec 2022

Bajaj Finserv Finance | India

- Developed Azure Data Factory ingestion pipelines integrating structured and semi-structured financial datasets into centralized analytics storage environments supporting enterprise reporting workloads.
- Engineered Spark transformation jobs in Azure Databricks processing high-volume financial transaction datasets, improving ETL processing throughput by nearly 30% across daily pipelines.
- Implemented Kafka streaming pipelines enabling near real-time availability of operational financial metrics used in Business Intelligence monitoring dashboards.
- Migrated legacy warehouse datasets into Synapse Analytics and ADLS lakehouse storage environments, contributing to infrastructure cost reduction of approximately 35% during platform transition.

- Automated SQL and Python validation frameworks detecting dataset inconsistencies across ingestion pipelines, improving downstream analytics reliability by 45%.
- Collaborated with reporting stakeholders to refine dimensional models supporting finance and operations dashboards used across multiple business units.
- Integrated anomaly-detection logic within Spark transformation pipelines identifying irregular transaction behavior earlier in processing workflows for analytics monitoring teams.
- Containerized repeatable transformation workloads using Docker to improve release consistency across development and staging pipeline deployment environments.

PROJECTS

Customer Behavior Analytics Pipeline

- Constructed a batch analytics pipeline using Apache Airflow and Amazon Redshift to aggregate multi-source customer interaction data, improving weekly Business Intelligence reporting readiness.
- Modeled dimensional reporting tables in Amazon Redshift using SQL transformations, improving campaign-performance visibility across segmentation dashboards for analytics teams.
- Validated dataset completeness using structured SQL reconciliation checks across ingestion layers, reducing reporting inconsistencies detected during scheduled analytics refresh cycles by 20%.

Streaming Operations Monitoring Data Platform

- Developed a Spark Streaming pipeline processing structured application event data, enabling near real-time service usage visibility across monitoring dashboards.
- Integrated pipeline outputs into Azure Synapse staging layers using Python transformations, improving downstream analytics dataset accessibility for reporting teams.
- Implemented pipeline health tracking using Azure Monitor alerts and logging signals, reducing issue detection time across streaming workflows supporting monitoring environments by 25%.

EDUCATION

Master of Science in Computer Science

University of Illinois Springfield | Springfield, IL

Bachelor of Computer Science

VelTech University | India

CERTIFICATIONS

- IBM Data Engineering Professional Certificate - **Coursera**
- Google Cloud Data Engineering Professional Certificate - **Coursera**
- Data Warehousing for Business Intelligence Specialization - **Coursera**
- Modern Big Data Analysis with SQL Specialization - **Coursera**
- Databricks Data Engineering Professional Certificate - **Coursera**
- Microsoft Azure Data Engineering Associate (DP-203) Preparation Specialization - **Coursera**